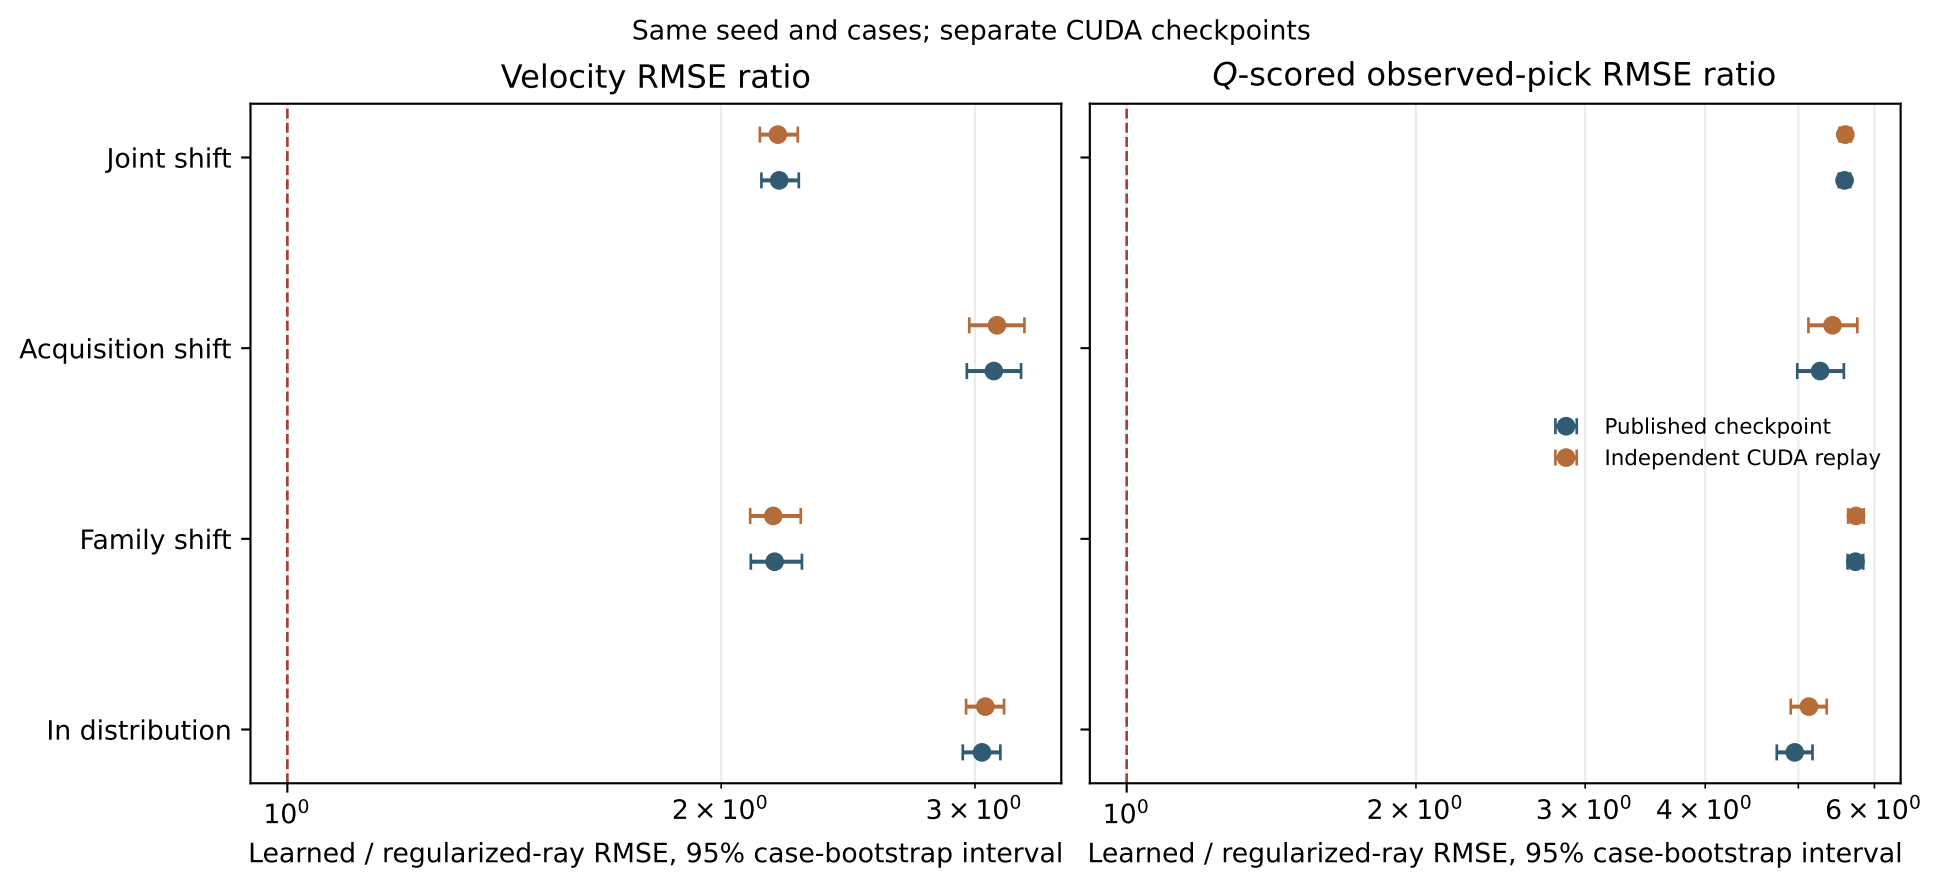
Same seed and cases; separate CUDA checkpoints
Velocity RMSE ratio
Q-scored observed-pick RMSE ratio
Joint shift
Acquisition shift
Family shift
In distribution
Published checkpoint
Independent CUDA replay
10^0
2 × 10^0
3 × 10^0
4 × 10^0
6 × 10^0
Learned / regularized-ray RMSE, 95% case-bootstrap interval
Learned / regularized-ray RMSE, 95% case-bootstrap interval